

# Letterboxd Recommendation System

---

Laya Gollamudi, Sierra Martinez-Kratz and Geraldine Nina Montano

APMA 4903: Senior Seminar

# Table of contents

1. About us
2. Problem Motivation
  - Letterboxd
  - Current recommendation system
3. Non-Negative Matrix Factorization (NMF)
4. Bayesian Personalized Ranking
5. Model Performance Comparison
6. Review Sentiment Analysis
7. Future work

# About us

---

# About us



Laya Gollamudi  
SEAS



Sierra Martinez-Kratz  
CC



Geraldine Nina  
SEAS

## Why we care?

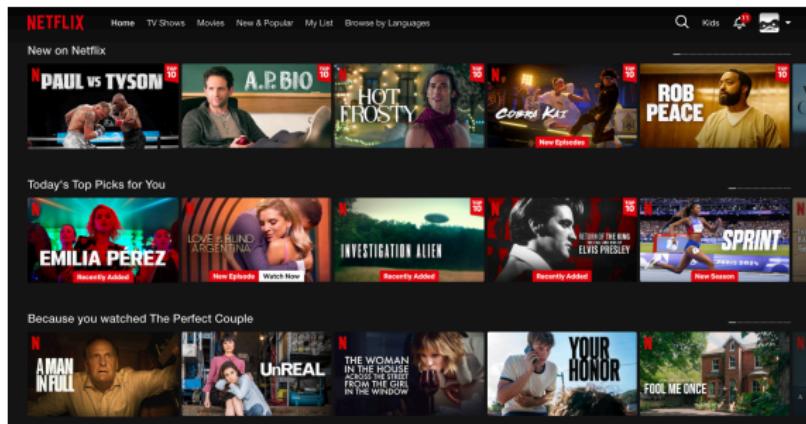
We are interested in—

- Data science
- Machine learning
- Letterboxd
- Movies!



# Why you should care?

- Recommendation systems suggest products, services, or content based on user preferences.
- Examples:
  - Movie recommendations (Netflix).
  - Product suggestions (Amazon).
  - Music streaming (Spotify).
- Understanding recommendation systems helps you understand why certain things are recommended to you



## References (1/3)

- Baad, D. (2020a). Sentiment Classification for Restaurant Reviews using TF-IDF. Retrieved from <https://medium.com/swlh/sentiment-classification-for-restaurant-reviews-using-tf-idf-42f707bfe44d>
- Burred, J. (2014). Detailed derivation of multiplicative update rules for NMF. [https://www.jjburred.com/research/pdf/jjburred\\_nmf\\_updates.pdf](https://www.jjburred.com/research/pdf/jjburred_nmf_updates.pdf)
- cheng, jianqiao. (2023, January 9). Understanding the Two-Tower Model in Personalized Recommendation Systems. Hackernoon.com. <https://hackernoon.com/understanding-the-two-tower-model-in-personalized-recommendation-systems>
- Colace, F., Conte, D., De Santo, M., et al. (2022a). A content-based recommendation approach based on singular value decomposition. \*Connection Science, 34\*(1), 2158–2176.  
<https://doi.org/10.1080/09540091.2022.2106943>
- Colace, F., Conte, D., De Santo, M., et al. (2022b). Recommender systems: A novel approach based on singular value decomposition. \*International Journal of Electrical and Computer Engineering (IJECE), 12\*(6), 6513–6521.  
<https://doi.org/10.11591/ijece.v12i6.pp6513-6521>
- Hug, N. (2020). Surprise: A Python library for recommender systems. \*Journal of Open Source Software, 5\*(52), 2174. <https://doi.org/10.21105/joss.02174>
- Kim, H., Park, H. (2008). Nonnegative Matrix Factorization Based on Alternating Nonnegativity Constrained Least Squares and Active Set Method. SIAM Journal on Matrix Analysis and Applications, 30(2), 713–730. <https://doi.org/10.1137/07069239x>
- Dang, C. N., Moreno-García, M. N., Prieta, F. D. I. (2021). An Approach to Integrating Sentiment Analysis into Recommender Systems. Sensors, 21(16), 5666.  
<https://doi.org/10.3390/s21165666>

## References (2/3)

---

- Learner, S. (n.d.). Letterboxd recommendations. GitHub. Retrieved November 17, 2024, from [https://github.com/sdl60660/letterboxd\\_recommendations/blob/main/README.md](https://github.com/sdl60660/letterboxd_recommendations/blob/main/README.md)
- Lee, D. D., & Seung, H. S. (2000). Algorithms for non-negative matrix factorization. \*Advances in Neural Information Processing Systems, 13\*, 556–562.
- Murray, I. (2022). Netflix Prize. Retrieved November 21, 2024, from University of Edinburgh Machine Learning and Pattern Recognition course notes:[https://mlpr.inf.ed.ac.uk/2022/notes/w8b\\_netflix\\_prize.pdf](https://mlpr.inf.ed.ac.uk/2022/notes/w8b_netflix_prize.pdf)
- Netflix. (2022). How Netflix's Recommendations System Works. Netflix Help Center.<https://help.netflix.com/en/node/100639>
- Pinoli, P., et al. (2015). Computational algorithms to predict gene ontology annotations. \*BMC Bioinformatics, 16\*(S6), Article S4. <https://doi.org/10.1186/1471-2105-16-s6-s4>
- Pu, L., & Faltings, B. (2013). Understanding and improving relational matrix factorization in recommender systems. \*Proceedings of the 2013 ACM International Conference on Recommender Systems\* (pp. 41–48). <https://doi.org/10.1145/2507157.2507178>
- Rajput, S., Mehta, N., Singh, A., Hulikal Keshavan, R., Vu, T., Heldt, L., ... Sathiamoorthy, M. (2023). Recommender systems with generative retrieval. Advances in Neural Information Processing Systems, 36, 10299–10315.
- Rendle, S., Freudenthaler, C., Gantner, Z., & Schmidt-Thieme, L. (2009). BPR: Bayesian personalized ranking from implicit feedback. \*Proceedings of the 25th Conference on Uncertainty in Artificial Intelligence (UAI 2009)\*, 452–461.

## References (3/3)

- Wankhade, M., Rao, A.C.S. & Kulkarni, C. (2022). A survey on sentiment analysis methods, applications, and challenges. \*Artificial Intelligence Review, 55\*, 5731–5780.  
<https://doi.org/10.1007/s10462-022-10144-1>
- How Tiktok recommends content. (n.d.). Retrieved from <https://support.tiktok.com/en/using-tiktok/exploring-videos/how-tiktok-recommends-content>
- Barnard, J. (2024). What are word embeddings? Retrieved from  
<https://www.ibm.com/topics/word-embeddings#%3A~\protect\protect\leavevmode@ifvmode\kern+.2222em\relaxtext=Word%20embeddings%20are%20a%20way, relationships%20among%20the%20corresponding%20words.>
- Recommenders Team. (n.d.). Cornac BPR deep dive. GitHub. Retrieved [Insert Date], from [https://github.com/recommenders-team/recommenders/blob/main/examples/02\\_model\\_collaborative\\_filtering/cornac\\_bpr\\_deep\\_dive.ipynb](https://github.com/recommenders-team/recommenders/blob/main/examples/02_model_collaborative_filtering/cornac_bpr_deep_dive.ipynb)
- E.M. Voorhees (1999). "Proceedings of the 8th Text Retrieval Conference" (PDF). TREC-8 Question Answering Track Report. pp. 77.
- Yining Wang, Liwei Wang, Yuanzhi Li, Di He, Wei Chen, Tie-Yan Liu. 2013. A Theoretical Analysis of Normalized Discounted Cumulative Gain (NDCG) Ranking Measures. In Proceedings of the 26th Annual Conference on Learning Theory (COLT 2013). pp 5.
- Zhai, J., Liao, L., Liu, X., Wang, Y., Li, R., Cao, X., ... Shi, Y. (2024). Actions speak louder than words: Trillion-parameter sequential transducers for generative recommendations. arXiv preprint arXiv:2402.17152.

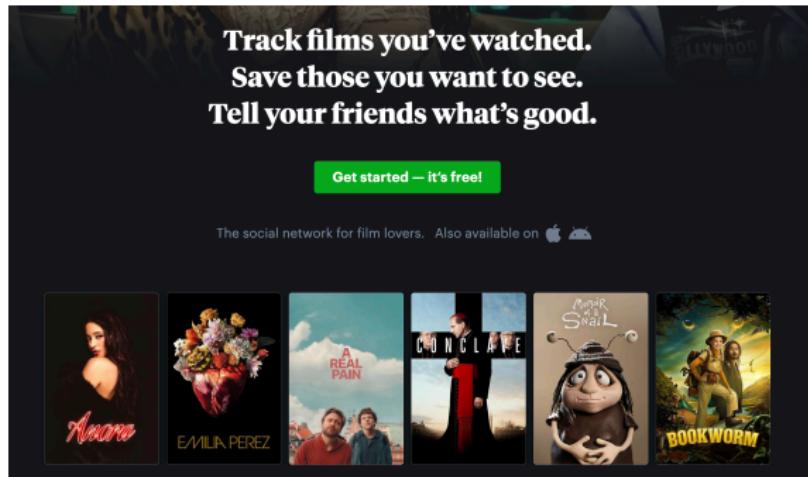
## Problem Motivation

---

# What is Letterboxd?

Like Beli, but for movies!

Founded in 2011



[letterboxd.com](https://letterboxd.com)

# How does it work?



**The Substance** 2024

 jer    42

Barbie for people who listen to Charli xcx

 14,737 likes

---



**Terrifier 3** 2024

 kani   22

shaking my head everytime art kills someone so the crowd at the theatre knows im against murder

 6,060 likes

---



**Challengers** 2024

 demi adejuyigbe   13

came home to my houseguest watching this for the first time and sat my ass down for the remainder. truly a shame this isn't message-driven enough to get love at the Oscars (I'm predicting best original screenplay and best score noms, maybe) because i do think this is one of the best films ever made about sex and power and attraction. love how each character gets a section to be the villain until they all feel equally villainous (and therefore,... [more](#)

 2,067 likes

# What's the problem?

No recommendation system :(

## People also ask :

Does Letterboxd give you movie recommendations?



Can Letterboxd generate recommendations for me? While we view the whole of Letterboxd as one big, organic recommendation engine, **we don't have a section specifically labeled as such**. You can, however, use our sorting and filters to find films that may be of interest.



Letterboxd

<https://letterboxd.com> › about › faq

[Frequent questions • Letterboxd](#)

# Sam Lerner's Solution

Created in December, 2020

Data was being automatically updated from Jan 2023-Jun 2024  
Hasn't pushed anything since a Google tag in September 2024

## Letterboxd Movie Recommendations

This site will gather movie ratings from any Letterboxd user and provide movie recommendations based on ratings data from thousands of other users.

The more movies you've rated on Letterboxd, the better and more personalized the recommendations will be, but it can provide recommendations to any user. If you'd like to filter out more well-known movies (based on the number of Letterboxd reviews), you can do so using the last slider below.

Letterboxd Username \*

Please provide a valid Letterboxd username

Faster Results —————— Better Results

All Movies —————— Less-Reviewed Movies Only

Add your ratings to recommendations database?

**GET RECOMMENDATIONS**

Source: [Sam Lerner's Letterboxd Recommendation Website](#)

# Sam Learner's solution

1.		Carol (2015)	Romance	Drama	10.00	
2.		Hamilton (2020)	History	Drama	10.00	
3.		Metropolis (1927)	Drama	Science Fiction	10.00	
4.		2001: A Space Odyssey (1968)	Science Fiction	Mystery	Adventure	10.00

How does it work?

Recommendations generated for Geraldine

Not great recommendations, mainly popular movies

- Problem: Predict user preferences
- Collaborative filtering is a popular approach for building recommendation systems.
- It relies on the behavior and preferences of other users to make predictions.
- Two main types:
  - **User-based collaborative filtering:** Recommends items based on similar users.
  - **Item-based collaborative filtering:** Recommends items based on the similarity between items.

Source: [\(Colace et al., 2022a\)](#)

# The User-Item Matrix

- Represents the relationship between users and items
- Users  $U = \{u_1, u_2, \dots, u_m\}$  and Items  $I = \{i_1, i_2, \dots, i_n\}$ .
- The **user-item matrix**  $R$  contains user preferences:

$$R_{ij} = \begin{cases} \text{rating given by user } u_i \text{ to item } i_j, & \text{if rated,} \\ \text{null,} & \text{otherwise.} \end{cases}$$

- Matrix  $R$  is usually sparse and quite large (most entries are missing)

Source: [\(Colace et al., 2022a\)](#)

# Simon Funk and SVD in the Netflix Prize

- **Context:**
  - Netflix launched a competition in 2006 to improve its recommendation engine by 10%.
  - The challenge involved predicting ratings in a large, sparse user-movie matrix.
- **Funk's SVD Approach:**
  - Simon Funk introduced an SVD method using stochastic gradient descent.
- **Impact and Popularization:**
  - Funk's approach quickly outperformed Netflix's existing system.
  - His open sharing of the method popularized SVD in collaborative filtering. [\[Funk's blog from 2006\]](#)
- **Legacy:**
  - Inspired further research and development in using SVD for recommendation systems.

Source: [\(Murray, 2022\)](#)

# Using SVD for Matrix Factorization

- SVD: Singular Value Decomposition
- SVD decomposes the user-item matrix  $R$ :

$$R \approx U\Sigma V^T$$

- $U$ : User-feature matrix (latent factors for users).
- $V$ : Item-feature matrix (latent factors for items).
- $\Sigma$ : Diagonal matrix with singular values (importance of each latent feature).

Source: [\(Colace et al., 2022b\)](#)

# Latent Factor Model

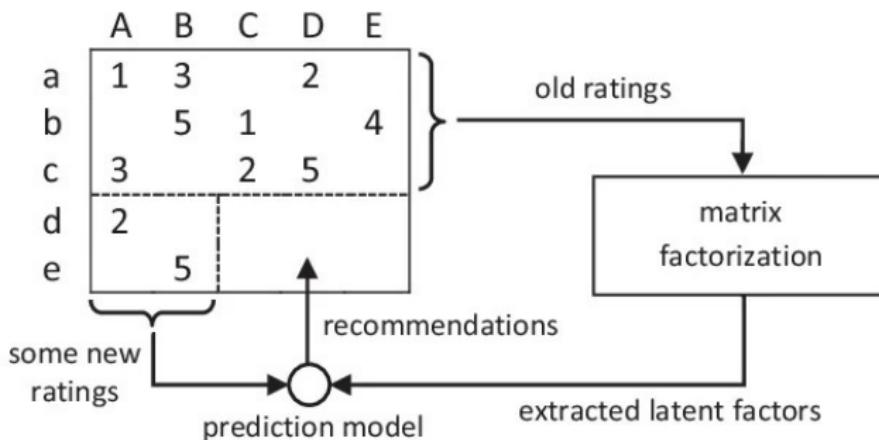
- Latent factors capture the underlying characteristics:
  - For movies: genre, director, or popularity.
  - For users: preferences for specific genres.
- Each user and item are represented by a vector in the latent space.
- Prediction: The user  $u_i$  rating for item  $j$  is:

$$\hat{R}_{ij} = \sum_{k=1}^K U_{ik} \Sigma_k V_{jk}^T$$

Sources: [\(Colace et al., 2022b\)](#), [\(Pu & Faltings, 2013\)](#)

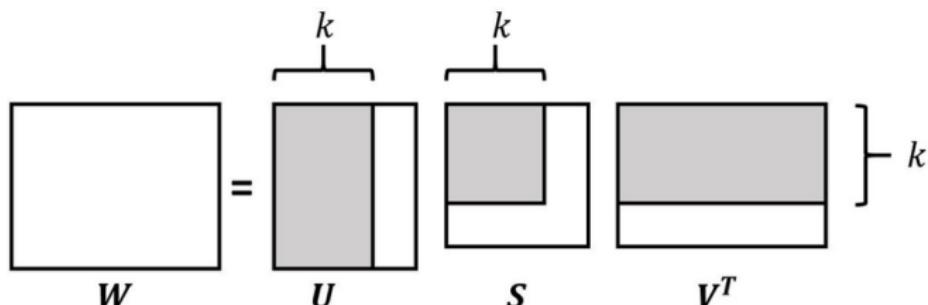
# Matrix Factorization Visualized

- "Old ratings" are fed into matrix factorization.
- Matrix factorization extracts latent factors, which represent hidden user and item characteristics.
- A prediction model combines these latent factors with any new ratings to make personalized recommendations.



# Truncated SVD for Recommendation Systems

- In practice, we use **truncated SVD** to retain only the top  $k$  singular values.
- Reduces complexity, dimensionality, and computational cost.
- Keeps only the most important latent factors without losing too much information.
- Simpler models tend to generalize better and are less likely to overfit, making more accurate predictions on unseen data.



Source: [\(Pinoli et al., 2015\)](#)

# Benefits of SVD for Recommendations

- Captures complex relationships between users and items.
- Works well for collaborative filtering—recommending items based on patterns across users.
  - Finds patterns across users to recommend items they are likely to enjoy.
  - Makes accurate predictions by leveraging these latent patterns.
- Truncated SVD reduces dimensionality, making it easier to handle large datasets with many users and items.
  - Focuses on the most meaningful factors, improving generalization.
  - Helps lower generalization error, leading to better performance on unseen data.

Sources: [\(Colace et al., 2022a\)](#), [\(Pinoli et al., 2015\)](#), [\(Pu & Faltings, 2013\)](#)

# Limitations of SVD for Recommendation Systems

- Often requires retraining when new data is added.
- Struggles with very sparse datasets and cold start problem.
- Assumes linear interactions between latent factors.
- Typically uses only explicit feedback.

Sources: [\(Colace et al., 2022a\)](#), [\(Pu & Faltings, 2013\)](#)

## Summary of SVD for Recommendations

- SVD factorizes the user-item matrix into latent user and item vectors, capturing hidden patterns.
- Latent factor model is used to predict missing ratings by identifying relationships between users and items, helping the model generalize
- SVD captures complex user-item relationships but has limitations, including struggles with sparse data and the cold start problem.

# Non-Negative Matrix Factorization (NMF)

---

# Non-Negative Matrix Factorization (NMF)

- **Definition:** NMF is a matrix factorization technique that decomposes a non-negative matrix  $R$  into two non-negative matrices  $W$  and  $H$  such that:

$$R \approx W \times H$$

- **Purpose:** Extracts latent features or patterns while preserving non-negativity, making results more interpretable in areas like topic modeling, image processing, and recommendation systems.

Source: ([Lee & Seung, 2001](#))

# NMF Problem Formulation

- **Objective:** Minimize the reconstruction error:

$$\min_{W,H} \|R - WH\|_F^2$$

- **Constraints:** Elements of  $W$  and  $H$  must be non-negative.
- **Interpretation:**
  - $W$ : Latent feature representation for rows of  $R$ .
  - $H$ : Latent feature representation for columns of  $R$ .

Source: [\(Lee & Seung, 2001\)](#)

# Frobenius Norm in Matrix Factorization

**Definition:** The Frobenius norm of the matrix  $R - WH$  is given by:

$$\|R - WH\|_F = \sqrt{\sum_{i=1}^m \sum_{j=1}^n |r_{ij} - (w_i^T h_j)|^2}$$

Where:

- $R \in \mathbb{R}^{m \times n}$  is the original matrix (e.g., user-item interactions).
- $W \in \mathbb{R}^{m \times k}$  and  $H \in \mathbb{R}^{k \times n}$  are the factorized matrices.
- $w_i^T h_j$  is the dot product between the  $i$ -th row of  $W$  and the  $j$ -th column of  $H$ .

## Key Properties:

- Measures the error or approximation between the original matrix  $R$  and the approximation  $WH$ .
- It is equivalent to summing the squared differences between the entries of  $R$  and  $WH$ .

# Comparison to Singular Value Decomposition (SVD)

- SVD: Factorizes a matrix  $R$  into  $U\Sigma V^T$ , where:
  - $U$  and  $V$  contain both positive and negative entries.
  - $\Sigma$  contains singular values representing the strength of each latent dimension.
- NMF also uses matrix factorization on explicit feedback
- Key Differences:
  - Non-negativity: NMF ensures that  $W$  and  $H$  are non-negative, unlike SVD.
  - Interpretability: NMF can yield more interpretable features since it preserves natural constraints.

Source: ([Lee & Seung, 2001](#))

# Solving NMF using Alternating Least Squares (ALS)

- **Basic Idea:** Fix one matrix (e.g.,  $W$ ) and solve for the other (e.g.,  $H$ ) alternately.
- **Optimization:** Minimizes the objective function by iteratively updating  $W$  and  $H$ .
- **Benefits of ALS:** Simplifies optimization by breaking it into convex subproblems.

Source: [\(Kim & Park, 2008\)](#)

# Pseudocode for NMF with Alternating Least Squares (ALS)

---

## Algorithm 1 NMF using ALS

---

- 1: Initialize  $W$  and  $H$  with non-negative values
- 2: **for** iteration in 1 to max\_iter **do**
- 3:   Fix  $H$ , update  $W$  by solving non-negative least squares
- 4:   **for** each row  $i$  of  $R$  **do**
- 5:      $W[i, :] \leftarrow \arg \min_{W[i,:]} \|H^T W[i, :] - R[i, :]\|_2^2$
- 6:   **end for**
- 7:   Fix  $W$ , update  $H$  by solving non-negative least squares
- 8:   **for** each column  $j$  of  $R$  **do**
- 9:      $H[:, j] \leftarrow \arg \min_{H[:,j]} \|WH[:, j] - R[:, j]\|_2^2$
- 10:   **end for**
- 11:   Compute reconstruction error; if below tolerance, break
- 12: **end for**
- 13: **return**  $W, H$

---

## Summary

- **NMF:** Powerful technique for non-negative data factorization, often yielding interpretable results.
- **Comparison to SVD:** NMF's non-negativity constraint aligns well with data like text, images and **reviews!**
- **ALS for NMF:** Alternates between  $W$  and  $H$  updates, providing a simple and effective solution.

# NMF Demo

(Code Demo)

# Bayesian Personalized Ranking

---

# Introduction to Bayesian Personalized Ranking (BPR)

- BPR is a ranking approach designed for recommendation systems with **implicit feedback**.
- Focuses on **pairwise ranking** of items based on observed interactions.
- Goal: Rank items a user has interacted with higher than items they haven't.

Source: ([Rendle et. al, 2009](#))

## Problem Setup

- Given a user  $u$  and items  $i$  and  $j$ :
  - $i$ : an item the user has interacted with (positive).
  - $j$ : an item the user has not interacted with (negative).
- Objective: Maximize the probability that the model ranks item  $i$  higher than  $j$ .

Source: [\(Rendle et. al, 2009\)](#)

## Score Calculation

- We are looking to decompose the user-item matrix into two matrices.

$$R \approx P \times Q^T$$

- Each user  $u$  and item  $i$  have associated **latent vectors**:

$$P_u \quad \text{and} \quad Q_i$$

- Predicted score for user  $u$  on item  $i$ :

$$\hat{x}_{ui} = P_u \cdot Q_i$$

- Goal: For positive item  $i$  and negative item  $j$ ,

$$\hat{x}_{ui} > \hat{x}_{uj}$$

Source: [\(Rendle et. al, 2009\)](#)

## Pairwise Preference Probability

- Define a probabilistic model for preference:

$$P(u, i, j) = \sigma(\hat{x}_{ui} - \hat{x}_{uj}) = \sigma(\hat{x}_{uij})$$

- Here,  $\sigma(x) = \frac{1}{1+e^{-x}}$  is the **sigmoid function**.
- This models the likelihood that  $u$  prefers  $i$  over  $j$ .

Source: [\(Rendle et. al, 2009\)](#)

# Log-Likelihood Objective

- Maximize the **log-likelihood** of observed preferences:

$$\log P(u, i, j) = \log \sigma(\hat{x}_{uij})$$

- Add a **regularization term** to prevent overfitting:

$$\log \sigma(\hat{x}_{uij}) - \lambda(\|P_u\|^2 + \|Q_i\|^2 + \|Q_j\|^2)$$

- Here,  $\lambda$  controls the regularization strength.

Source: [\(Rendle et. al, 2009\)](#)

# Optimization with Gradient Descent

- Use **gradient descent** to update parameters:
  - User vector  $P_u$
  - Item vectors  $Q_i$  and  $Q_j$
- Adjust vectors to maximize the likelihood of positive item scores over negative item scores.

Prediction Formula:

$$\hat{x}_{ui} = P_u \cdot Q_i = \sum_{f=1}^k P_{uf} \cdot Q_{if}$$

Derivatives of Matrix Factorization:

$$\frac{\partial}{\partial \theta} \hat{x}_{uij} = \begin{cases} (Q_{if} - Q_{jf}) & \text{if } \theta = P_{uf} \\ P_{uf} & \text{if } \theta = Q_{if} \\ -P_{uf} & \text{if } \theta = Q_{jf} \end{cases}$$

## Pseudocode for BPR (Part 1: Initialization)

---

### Algorithm 2 CREATEBPRTriplets( $R$ )

---

```
1: Input: User-item interaction matrix  $R$ 
2: Output: Set of training triplets  $D_R$ 
3:  $D_R \leftarrow \emptyset$  {Initialize the set of training triplets}
4: for each user  $u \in U$  do
5:   for each item  $i \in I$  such that  $(u, i) \in R$  do
6:     for each item  $j \in I$  such that  $(u, j) \notin R$  do
7:        $D_R \leftarrow D_R \cup \{(u, i, j)\}$  {Add triplet to  $D_R$ }
8:     end for
9:   end for
10: end for
11: return  $D_R$ 
```

---

Source: ([Rendle et. al, 2009](#))

## Pseudocode for BPR (Part 2: Training)

---

**Algorithm 3** LEARNBPR( $D_R, \Theta$ )

---

- 1: Randomly initialize  $P$  and  $Q$
  - 2: **repeat**
  - 3:   Draw  $(u, i, j)$  from  $D_R$
  - 4:    $P_u \leftarrow P_u + \eta ((1 - \sigma(\hat{x}_{uij}))(Q_i - Q_j) - \lambda P_u)$
  - 5:    $Q_i \leftarrow Q_i + \eta ((1 - \sigma(\hat{x}_{uij}))P_u - \lambda Q_i)$
  - 6:    $Q_j \leftarrow Q_j + \eta (-(1 - \sigma(\hat{x}_{uij}))P_u - \lambda Q_j)$
  - 7: **until** convergence
  - 8: **return**  $P, Q$
- 

Source: ([Rendle et. al, 2009](#))

# Model Performance Comparison

---

## Precision@k

**Definition:** Precision@k measures the fraction of relevant items in the top- $k$  recommendations.

**Formula:**

$$\text{Precision}@k = \frac{\text{Number of relevant items in top-}k}{k}.$$

**Example:** For recommendations {A, B, C} with relevant items {C, D}:

- Top- $k$ : {A, B, C},  $k = 3$ .
- Precision@3 =  $\frac{1}{3}$ .

# Reciprocal Rank (RR)

**Definition:** RR is reciprocal rank of the first relevant item in a ranked list.

**Formula:**

$$RR = \frac{1}{\text{rank}_i},$$

where  $\text{rank}_i$  is the position of the first relevant item in user  $i$ 's recommendation list.

**Example:** For recommendations {A, B, C} with relevant item {C}:

- Rank of C is 3.
- $RR = \frac{1}{3}$ .

Source: ([Voorhees, 2007](#))

# Normalized Discounted Cumulative Gain (NDCG)

**Definition:** NDCG measures the quality of ranking by considering the positions of relevant items, assigning higher weights to items ranked higher.

**Formula:**

$$\text{NDCG} = \frac{\text{DCG}}{\text{IDCG}},$$

where

$$\text{DCG} = \sum_{i=1}^k \frac{\text{rel}_i}{\log_2(i+1)},$$

and IDCG is the DCG of the ideal ranking.

**Example:** Recommendations: {A, B, C}, Relevant items: {A: 3, C: 2, D: 1}.

- $\text{DCG} = \frac{3}{\log_2(1+1)} + \frac{0}{\log_2(2+1)} + \frac{2}{\log_2(3+1)} = 3.5.$
- $\text{IDCG} = \frac{3}{\log_2(1+1)} + \frac{2}{\log_2(2+1)} + \frac{1}{\log_2(3+1)} = 4.76.$
- $\text{NDCG} = \frac{3.5}{4.76} \approx 0.735.$

# General and user-level performance comparison

Code demo

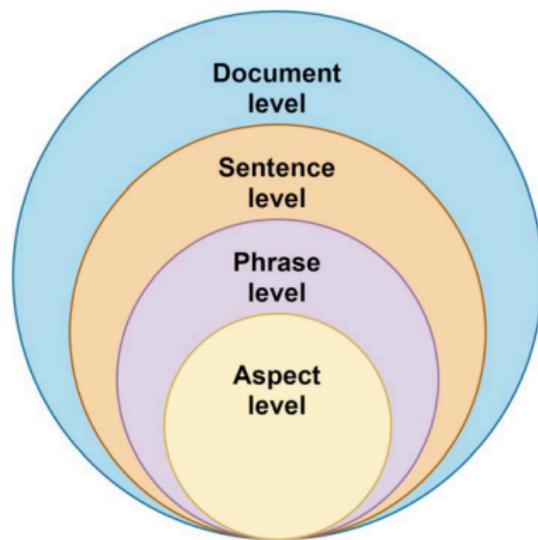
# Review Sentiment Analysis

---

# Introduction to Sentiment Analysis

- Extracts subjective information from text using NLP and text mining
- Common applications
  - Analyze customer feedback
  - Identify/track market trends
  - Monitor brand reputation

# Levels of Sentiment Analysis



Source: [A survey on sentiment analysis methods, applications, and challenges](#)

# Levels of Sentiment Analysis

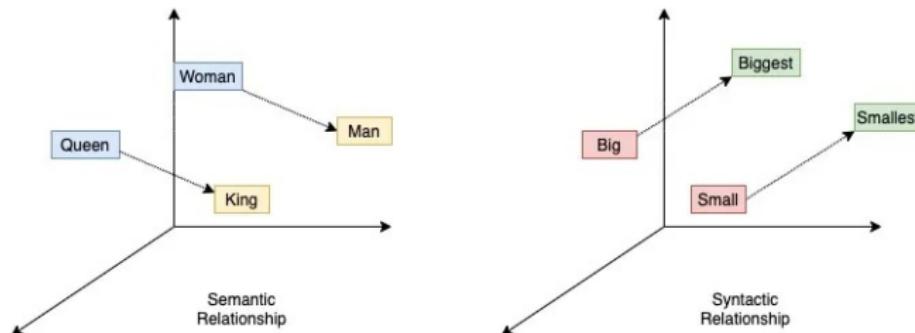
- Document level
  - For example, classify chapters of a book
  - Not as common
- Sentence level
  - Useful when a document has a wide range of sentiments
  - Subjective classification
  - Classified sentiments can be used individually or aggregated to classify the sentiment of the entire document
- Phrase level
  - Useful for product reviews that contain multiple lines
- Aspect level
  - Sentiment classification of individual words
  - Classify sentiment of each aspect and aggregate to classify sentiment of entire sentence

## Feature Extraction

- Extract most useful information that represents the text's most important features
- Terms frequency: number of times a word has appeared
- Parts of speech tagging: Classify token as its part of speech (i.e. noun, verb, adjective, etc.)
- Negation: words that can reverse polarity (e.g. not, never, none, etc.)
- Bag of words: combine vocabulary of words into a "bag"
  - Problem: doesn't consider syntactical meaning

# Word Embedding

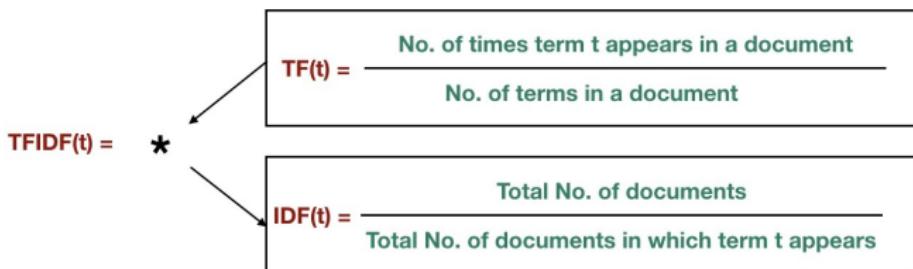
- Represent words in a vector space by grouping words with similar meanings
- Distance and direction between vectors are determined by the similarity and relationships between words



Source: [Word2Vec Research Paper Explained](#)

# TF-IDF

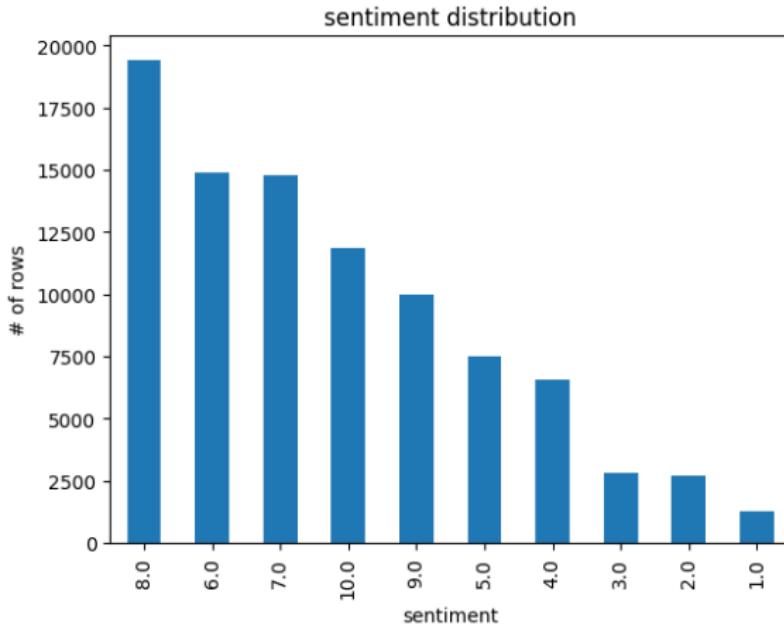
Frequency-based embedding that assigns a weight to each word that represents the significance of the word (first introduced in 1972)



Source: [Sentiment Classification for Restaurant Reviews using TF-IDF](#)

# Outline of Sentiment Analysis - Balance Classes

**Problem:** imbalanced classes



# Outline of Sentiment Analysis - Tokenization

- Split text into an array of words (tokens)
- Allows for each word to be transformed into a number

tokens	
4	[if, the, real, pompeii, was, this, boring, gl...]
16	[eli, roth, needs, to, stick, to, horror, even...]
50	[should, ve, stayed, at, hilton, hotel]
100	[what, pretentious, pile, of, horse, dung]
141	[would, like, to, formally, apologize, to, the...]

# Outline of Sentiment Analysis - Stemming

- Remove prefixes and suffixes to get the root word of a word
- Choose between PorterStemmer, LancasterStemmer, and SnowballStemmer based on accuracy of results
- We choose **SnowballStemmer** (developed 2013-2019)

porter_stemmed	
4	[if, the, real, pompeii, wa, thi, bore, glad, ...]
16	[eli, roth, need, to, stick, to, horror, even,...]
50	[should, ve, stai, at, hilton, hotel]
100	[what, pretenti, pile, of, hors, dung]
141	[would, like, to, formal, apolog, to, the, nat...]

lancaster_stemmed	
4	[if, the, real, pompeii, was, thi, bor, glad, t...]
16	[eli, roth, nee, to, stick, to, hor, ev, though...]
50	[should, ve, stay, at, hilton, hotel]
100	[what, pretenty, pil, of, hors, dung]
141	[would, lik, to, form, apolog, to, the, nat...]

snowball_stemmed	
4	[if, the, real, pompeii, was, this, bore, glad...]
16	[eli, roth, need, to, stick, to, horror, even,...]
50	[should, ve, stay, at, hilton, hotel]
100	[what, pretenti, pile, of, hors, dung]
141	[would, like, to, formal, apolog, to, the, nat...]

## Outline of Sentiment Analysis - TF-IDF

- Build dictionary from stemmed tokens
- Build corpus from converting stemmed tokens of each review into a bag of words
- Build TF-IDF model from corpus
- Generate TF-IDF vectors for each review
- Use DecisionTreeClassifier to perform sentiment classification

```
sample data from dictionary:  
word: bore - id: 0  
word: glad - id: 1  
word: got - id: 2  
word: if - id: 3
```

## Combined Recommendation Method

Given  $m$  users and  $n$  items in the training rating matrix  $R_{m \times n}$  ( $r_{ij} \in R_{m \times n}$ : rating of user  $u_i$  on item  $i_j$ ), the predicted rating of user  $u_a$  on item  $i_j$  in the test set is given by:

$$pr_{aj} = \beta * pr\_mf_{aj} + (1 - \beta) * pr\_sent_{aj}$$

- $pr\_mf_{aj}$ : predicted rating by matrix factorization methods for user  $u_a$  on item  $i_j$
- $pr\_sent_{aj}$ : predicted rating by sentiment model for user  $u_a$  on item  $i_j$
- $\beta$  adjusts the weight/importance of each term

# Rating Prediction from Sentiment Model Pseudocode

---

## Algorithm 4 Compute $pr\_sent_{aj}$

```
1: Initialize list of items  $I$  to an empty list
2: for item  $i_k$  in training set do
3:   if user  $u_a$  has rated item  $i_k$  then
4:     Add  $i_k$  to  $I$ 
5:   end if
6: end for
7: Initialize list of users  $U$  to an empty list
8: for user  $u_b$  in training set do
9:   for item  $i_k$  in  $I$  do
10:    if user  $u_b$  has rated item  $i_j$  and user  $u_b$  has rated item  $i_k$  then
11:      Add user  $u_b$  to  $U$ 
12:    end if
13:   end for
14: end for
15: Initialize list of similarities  $S$  to an empty list
16: if  $U$  is not empty then
17:   for user  $u_i$  in  $U$  do
18:      $s_{ai} \leftarrow$ cosine_similarity(user  $u_a$ , user  $u_i$ ) {using sentiment model}
19:     Add  $s_{ai}$  to  $S$ 
20:   end for
21:   Calculate  $pr\_sent_{aj}$ 
22: end if
```

---

## Relevant Equations in Sentiment Model Pseudocode

$$s_{ai} = \frac{\sum_{j=1}^n r_{aj}r_{ij}}{\sqrt{\sum_{j=1}^n r_{aj}^2} \sqrt{\sum_{j=1}^n r_{ij}^2}} \text{ (cosine similarity)}$$

$$pr\_sent_{aj} = \bar{r}_a + \frac{\sum_{i=1}^k s_{ai}(r_{ij} - \bar{r}_i)}{\sum_{i=1}^k |s_{ai}|}$$

- $r_{aj}$ : rating of user  $u_a$  on item  $i_j$
- $r_{ij}$ : rating of user  $u_i$  on item  $i_j$
- $\bar{r}_a$ : average rating of user  $u_a$
- $\bar{r}_i$ : average rating of user  $u_i$
- $s_{ai}$ : similarity between user  $u_a$  and user  $u_i$

# Sentiment Analysis Results

Code demo

## Future work

---

## Future of the field: Netflix

- Considers similar factors like viewing history, ratings, others with similar preferences, and information about movies such as genre, cast, year released
- Takes into account time of day, user's preferred language, time constraints, and the device being used
- Personalizes the layout of movies, such as the row they appear in and the order of the titles, and user searches



Source: [Netflix](#)

## Future of the field: Amazon

- Considers item-to-item similarity instead of user-to-user similarity.
- Uses [multi-layered neural networks](#) for Prime Video recommendations. Considers watch history, time of day, day of week and even weather.



Source: [Amazon](#)

## Future of the field: Tiktok

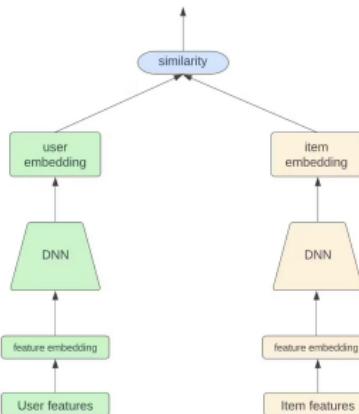
- Considers three main factors: user interactions, content information, user information
- Predictions also influenced by users with similar interests
- Example: For You
  - User interactions: likes, shares, comments, length of time spent watching a video
  - Content information: sounds, hashtags, number of views, country in which video was published
  - User information: device settings, language preference, location, time zone, day, device type



Source: [TikTok](#)

# Future of the Field, More Generally

- Generative recommenders (2024): predict sequences of user behavior, well-suited for large amounts of frequently-changing data
- Generative retrieval (2023): predicts items a user is likely to interact with in the future, well-suited for handling new items
- Two-tower model:



Source: Understanding the Two-Tower Model in Personalized Recommendation Systems

# Our futures



**Laya Gollamudi**  
SEAS  
Management  
Consulting



**Sierra Martinez-Kratz**  
CC  
Data Science  
Grad School



**Geraldine Nina**  
SEAS  
Risk Analysis

Questions?